



Everything you wanted to know, but were too afraid to ask about...

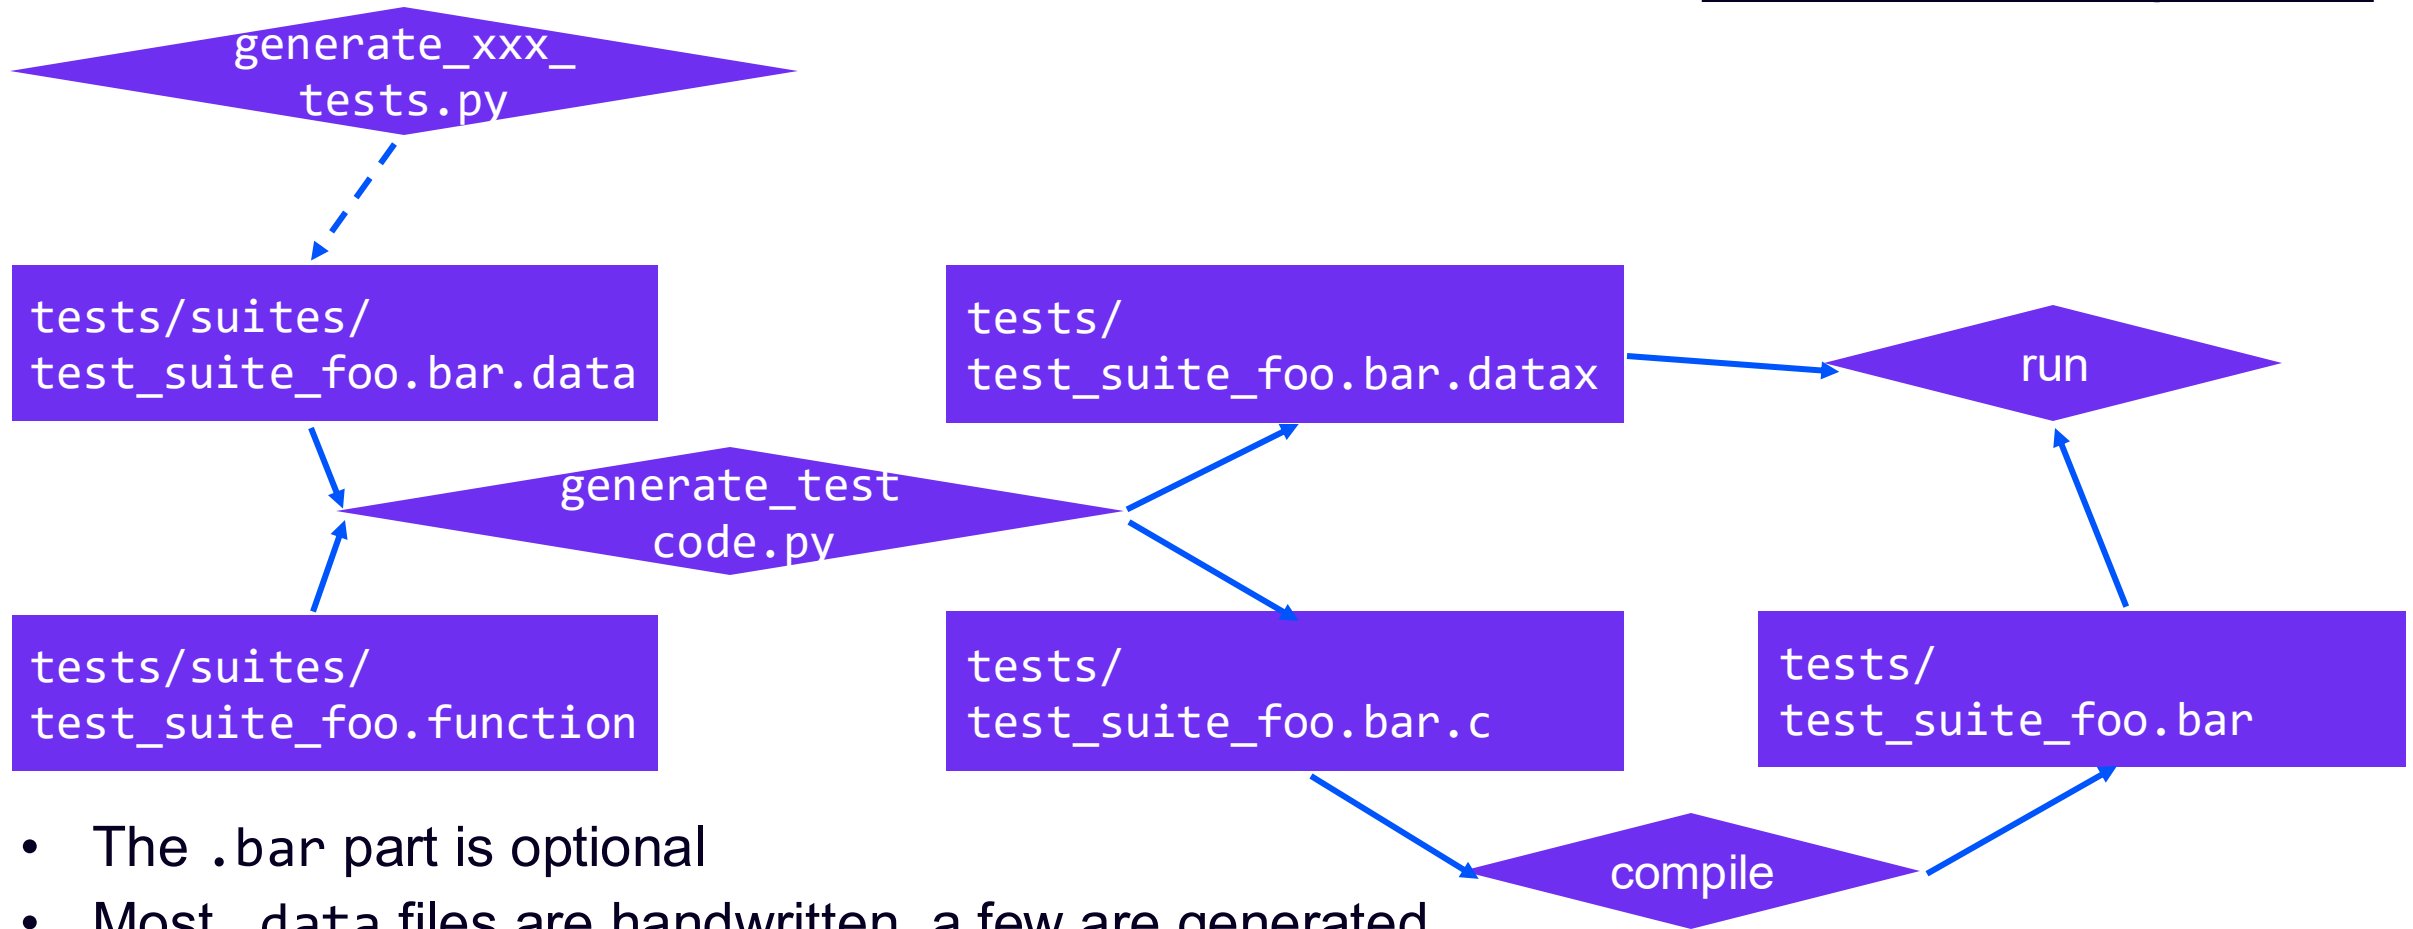
# Mbed TLS and TF-PSA-Crypto unit tests

Gilles Peskine and others  
2025-07-04

# Writing unit tests

# Compilation

Doc: [kb: Mbed TLS tests guidelines](#)



- The .bar part is optional
- Most .data files are handwritten, a few are generated
- No need to declare new files in build scripts
  - But you need to re-run cmake

# Test assertions

- Effects of `TEST_ASSERT(condition)`:
  - Mark test case as failed
  - `goto exit`
    - Not return because we often need to clean up
    - `generate_test_code.py` adds `exit;`; if there's no `exit` label
  - Display the failure location when the test exits
- Fancier assertions
  - `TEST_EQUAL(a, b)` asserts `a==b`
    - Displays the values on failure
    - Also `TEST_LE_S(a, b)` to assert `a<=b` with `a, b` signed
    - Also `TEST_LE_U(a, b)` to assert `a<=b` with `a, b` unsigned
  - `TEST_MEMORY_COMPARE(buffer1, length1, buffer2, length2)`
    - asserts the same length and content
  - `TEST_CALLOC(pointer_variable, size)`
    - Allocate memory (free it with `mbedtls_free()`)
    - `pointer_variable` must be null
    - The size is in elements, not bytes (e.g. `int *p = NULL; TEST_CALLOC(p, 3);` → 3 integers)
  - More: see <tests/include/test/macros.h>

# Limitations on information displayed on failure

- Failure information is displayed when the test case returns
  - Only a limited space to store information between mark-as-failed and test-returns
  - Static strings: function, file name, message
  - Integers: line, lhs, rhs, step
- Tip: some complex conditions can be restructured for better display

Basic	Nicer
<code>TEST_ASSERT(cond1 &amp;&amp; cond2);</code>	<code>TEST_ASSERT(cond1); TEST_ASSERT(cond2);</code>
<code>TEST_ASSERT(ret == 0                  ret == ACCEPTABLE_ERROR);</code>	<code>if (ret != 0) {     TEST_EQUAL(ret, ACCEPTABLE_ERROR); }</code>
<code>TEST_ASSERT(x &lt; max);</code>	<code>TEST_LE_S(x, max - 1);</code>

# Test step

- You only see where the failure happened, no stack trace and no local state
- `MBEDTLS_TEST_SET_STEP(n)`
  - Shows “at step n” in the failure message

## Example: loop

```
for (i = 0; i < count; i++) {  
    mbedtls_test_set_step(i);  
    do_stuff(i);  
    TEST_ASSERT(check(i));  
}
```

## Example: auxiliary function called several times

```
mbedtls_test_set_step(1);  
aux_may_fail(...);  
mbedtls_test_set_step(2);  
aux_may_fail(...);  
mbedtls_test_set_step(3);  
aux_may_fail(...);
```

- Limitation: does not nest
  - The step is a global variable
  - For nested loops, you can use e.g. `MBEDTLS_TEST_SET_STEP(outer_i * 100000 + inner_i)`

# Running unit tests

# Skip boring output

- To run all test cases but omit skip/pass lines:
  - `tests/test_suite_foo |& grep -Ev '(PASS|SKIP|----)'`
  - If you don't like huge test suites, split them!
- To skip slow test suites:
  - With Make, set the environment variable `SKIP_TEST_SUITES`, e.g.  
`SKIP_TEST_SUITES=constant_time_hmac,lmots,lms,gcm,psa_crypto.pbkdf2,ssl_decrypt make test`
  - With CMake, set the CMake parameter `SKIP_TEST_SUITES` (you have to re-run cmake to change it), e.g.  
`cmake -B build-debug -DCMAKE_BUILD_TYPE=Debug  
-DSKIP_TEST_SUITES=constant_time_hmac,lmots,lms,gcm,psa_crypto.pbkdf2,ssl_decrypt`



# Run only one test case

- *I'm debugging one particular test case. I set a debugger breakpoint but it's triggered by many previous test cases. How can I skip over them?*
- Hack: delete the previous cases in \*.data
  - Assumes you haven't modified them!
  - Remember to undelete before committing!
- Hack: comment out the previous cases in \*.data
  - Remember to uncomment before committing!
- Hack: copy the test case to the top of \*.data
  - If you change the test case, remember to update both copies!
  - Remember to remove the extra copy before committing!
- Hack: copy to a new file test\_suite\_foo.temp.data
  - If you change the test case, remember to update both copies!
  - make picks it up automatically, but cmake and mbedtls-prepare-build must be re-run
- Hack: copy the test case in \*.datax to a new file and run tests/test\_suite\_foo my.datax
  - Difficult to change the data

# Debugging unit tests

We use bullets on level one

Here is our second level bullet

Here is our next level bullet

Here is our next level bullet

# Which function failed first?

- *bedtls\_foo() returned -1234, how do I find which of the 50 functions it might have called detected the error?*
- Use a debugger that supports reverse stepping (a.k.a backward debugging a.k.a. time travel a.k.a...)
  - E.g. Visual Studio (Windows only) or gdb (Linux)
    - macOS: [warpspeed](#)?
  - Set a breakpoint after mbedtls\_foo() returns
  - Step back to the first return MBEDTLS\_ERR\_EXAMPLE
- *Gdb? Really? I never got it to work!*

# Reverse debugging made simple on Linux

- Initial setup
  1. Install the Mozilla Record and Replay framework (rr) <https://rr-project.org/> (apt install rr).
  2. If needed, give yourself debugging permission: `sudo sysctl kernel.perf_event_paranoid=1` (the Ubuntu default is 4 which is too paranoid).
  3. To make this survive across reboots: `echo 'kernel.perf_event_paranoid = 1' >>/etc/sysctl.d/zz-local.conf`
- Debug a program
  1. Build with debugging symbols as usual (`-O0 -g3` or `-Og -g3`).
  2. `rr record ./test_suite_ssl` saves a full trace of the execution.
  3. `rr replay` gives you a gdb interface where reverse execution actually works.
    - Use reverse-xxx commands
    - `rs` (reverse-step) steps into functions
    - `rn` (reverse-next) steps over function calls
    - `reverse-finish` goes back to where the current function was called
  4. If you use a frontend, configure it to run `rr replay` instead of `gdb myprogram`.
    - If the frontend uses gdb's machine interface: `rr replay -i=mi ...` instead of `gdb -i=mi ...`

# Meaning of numerical values

- *When `TEST_EQUAL(a, b)` fails, it shows numerical values. What about symbolic names for enum-like types?*
- Guess the type from the assertion and search the source code.
- For Mbed TLS error codes: `programs/util/strerror` (in `mbedtls`)  

```
$ mbedtls-strerror -28160
```

Last error was: -0x6e00 - SSL - The handshake negotiation failed
- If you want the macro name:  

```
$ git grep --recurse-submodules '#define MBEDTLS_ERR_.*-0x3F00'
```

`tf-psa-crypto/include/mbedtls/pk.h:#define MBEDTLS_ERR_PK_TYPE_MISMATCH -0x3F00`
- For PSA constants: `programs/psa/psa_constant_names`  

```
$ psa_constant_names alg 100664841
```

`PSA_ALG_ECDSA(PSA_ALG_SHA_256)`
- Tip: constants change rarely so you can install `psa_constant_names` and `strerror` from any version

# Debugging builds with PSA drivers

- Hackish: edit an `all.sh` component to get a build with debug symbols:
  - Add `ASAN_CFLAGS=' -O0 -g3 '` at the beginning
  - Replace `make test` by `false`
  - Run `all.sh` (without `-k`) and you'll get the drivers and library built with debug symbols
  - You can't do incremental builds. Tip: use [ccache](#) for faster rebuilds:  
`CC="ccache ${CC:cc}" ASAN_CC="ccache clang"`
- `mbedtlsls-prepare-build` has some code for driver builds but it's still unstable and clumsy
- ???

# SSL unit tests

- *If an SSL test doesn't enable debug logs, can I see them anyway when debugging? I'd like something as simple as adding `debug_level=4` to `ssl-opt.sh`.*
- As of June 2025:
  - Set up the test for debugging:  
`options.cli_log_fun = mbedtls_test_ssl_log_analyzer;`  
`options.srv_log_fun = mbedtls_test_ssl_log_analyzer;`
  - In `tests/src/test_helpers/ssl_helpers.c`, change `#if 0` to `#if 1`
  - All debug logs go to stdout
  - Run `tests/test_suite_ssl -v` (without `-v`, stdout is suppressed)
- Simpler (just set a variable and run with `-v`)
  - coming in <https://github.com/Mbed-TLS/mbedtls/pull/10273>

# Unit tests on the CI



# Which components have failed?

- If test cases have failed, this is recorded in `failures.csv`
  - Just the FAIL lines of `outcomes.csv`
  - Available under “Artifacts”
  - Tip: if the list on BlueOcean only shows `all*.log.gz`, copy one URL and edit it!
  - Note: the file is compressed if it's large, so try `failures.csv.xz`
- The outcome file only records test case failures, not e.g. build failures
  - Improvement pending: [Mbed-TLS/mbedtls-framework#129](#)
  - Pipeline Steps + userscript: see [CI failure FAQ “PR tests failed. Which components failed?”](#)
  - Anonymous tip *“it can be found by looking for “Failed jobs” in the pipeline log on the artefacts page, but I can never remember that “Failed jobs” is the string to look for - something more obvious would be helpful.”*
  - Tip from Bence: add `/wfapi` to the classic (non-BlueOcean) main page  
(`/job/mbed-tls-pr-head/job/PR-NNNN-head/1/wfapi`) for a JSON list of failed jobs

# Analyzing outcome files

- Outcome file format (documented in [test-framework.md](#))
  - platform;component[-configuration];test\_suite;test case description;FAIL;message
  - E.g. in which configurations did this test case fail and why?
    - `grep ';test_suite_foo;test case description;FAIL;' failures.csv | cut -d';' -f2,6`
- *Which test cases are failing, across all components?*
  - `<failures.csv cut -d\; -f3,4 | sort -u`
  - `xzcat failures.csv.xz | ...` if the file is compressed
  - `... | sort | uniq -c | sort -n` to sort by the count of failing components
- *A test case is failing only in some configurations. What's special about them?*
  - First, determine in which components the test case is passing vs failing.  
`xzcat outcomes.csv.xz | grep ';test_suite_foo;Description of the test case;PASS;' | cut -d\; -f2`  
`xzcat outcomes.csv.xz | grep ';test_suite_foo;Description of the test case;PASS;' | cut -d\; -f2`
  - Next you can try to analyze the configurations.
    - The SKIP/PASS status for `test_suite_config*` tells you about each component's configuration.
    - `framework/scripts/search_outcomes_config.py` queries the outcome file for config properties.

# Reproduce an all.sh build with debugging

- General tip: to see the products from an `all.sh` build, add the command `false` before running the tests and run `all.sh` without `-k`. It will stop on the failure without cleaning up.
  - `all.sh` cleans up on Ctrl+C. Use Ctrl+Z to suspend it. Use Ctrl+\ or `kill -9` to kill it without cleanup.
- If it's an ASan build using `make` (not `CMake`), add `ASAN_CFLAGS=' -Og -g3 '` at the beginning of the component, and replace `make test` by `false`
- With `CMake`, add/change the build type: `cmake -DCMAKE_BUILD_TYPE=Debug` or `cmake -DCMAKE_BUILD_TYPE=ASanDbg`
- ???

arm

Merci

Danke

Gracias

Grazie

谢谢

ありがとう

Asante

Thank You

감사합니다

धन्यवाद

Kiitos

شكراً

ধন্যবাদ

תודה

ధన్యవాదములు

Köszönöm